

Latency-aware end-to-end (composition): decentralized detection adds a degraded window before FL $\rightarrow$ GL, then GL recovers ($N = 14 \rightarrow 7$, $k \geq 2$)

