

OUTER — held-out, evaluated once

ERENO train

2,955,738 records

ERENO test

2,955,648
held out

INNER — all tuning on the training split

benign cap → 696,313 pool
per-client 80/20 (local train / validation)
GRASP feature selection (5-fold CV) → 24 features
 λ penalty → tuned on validation
k quorum, threshold $\tau = 0.5$ → reported levers
train N attack-specialist boosters

evaluate once

Single evaluation

F1 / recall / FPR / #FP
on the held-out test