



Convergence of a NEW attack: learning (vs data) THEN dissipation — injection vs random_replay

1 · LEARNING — training vs accumulated data (10 seeds, mean $\pm$ std)
(local, mode-independent: same for FL and GL)



2 · DISSIPATION — spreading the trained specialist
FL ~ 1 round; GL diffuses over $\sim \Theta(N)$ ($N = 12$, ring)



Both attacks are Tier-2-detectable (coherent autonomous discovery). Learning latency = collection rounds to a good specialist; dissipation = FL depth-1 vs GL $\sim \Theta(N)$.